

NIAPythonDay3

March 27, 2019

NIA Python Bootcamp UNIT 3 - Wednesday July 19, 2017

1 UNIT 1 review

1. Python ecosystem of tools
2. Jupyter Notebook is code, output and documentation all in one document
3. Type code into cells, and to run them you press Shift-Enter
4. Different data types for different data
5. Tab completion reduces typing, shows you pop-up menu of all the things you can do with that piece of data
6. Operators take one or more input values and turn them into other values *based on the input values type*
7. Converting data from one type to another using the function syntax, e.g., int()

2 UNIT 2 Review

1. Exploring data types using the TAB key
2. Python syntax for taking slices of iterables
3. NumPy arrays: basic math operations in 1-D and 2-D (e.g., row-wise and column-wise eman)
4. Subselecting based on a boolean criterion
5. Example: Images as 3-D matrices

3 UNIT 3:

3. PANDAS DataFrames
4. Simple and complex sorting

3.1 PANDAS DataFrame

- pandas = [Python Data Analysis Library](#)
- Emulate R's data.frame structure.
- Basically a NumPy matrix with
 - Row and column names
 - Can have columns of different types
 - Handles missing data better

3.2 Load the PANDAS package into memory using import()

```
In [1]: import pandas as pd
```

3.3 Use PANDAS read_* functions to import data

- There are many functions to import data
- Type `pd.read_` then TAB to see all the import functions

```
In [ ]: pd.read_
```

3.4 Read data from file or URL

```
In [2]: titanic_data_url = "http://biostat.mc.vanderbilt.edu/wiki/pub/Main/DataSets/titanic3.."
```

```
In [3]: titanic = pd.read_excel( titanic_data_url )
```

3.5 Return type is a DataFrame

```
In [4]: type(titanic)
```

```
Out[4]: pandas.core.frame.DataFrame
```

3.6 What did we just load?

```
In [5]: titanic.shape
```

```
Out[5]: (1309, 14)
```

3.6.1 Change the number of rows Pandas will display using the `set_option()` function

Use the word `None` if you want to display all of them.

```
In [6]: pd.set_option( 'display.max_rows', 50 )
```

3.6.2 See the first N rows using `.head(N)`

Defaults to first 5

```
In [7]: titanic.head(2)
```

```
Out[7]:
```

	pclass	survived	name	sex	age	sibsp	\
0	1	1	Allen, Miss. Elisabeth Walton	female	29.0000	0	
1	1	1	Allison, Master. Hudson Trevor	male	0.9167	1	

	parch	ticket	fare	cabin	embarked	boat	body	\
0	0	24160	211.3375	B5	S	2	NaN	
1	2	113781	151.5500	C22 C26	S	11	NaN	

	home.dest
0	St Louis, MO
1	Montreal, PQ / Chesterville, ON

3.6.3 See the last N rows using .tail(N)

Defaults to last 5.

```
In [8]: titanic.tail(1)
```

```
Out[8]:
```

	pclass	survived	name	sex	age	sibsp	parch	ticket	\
1308	3	0	Zimmerman, Mr. Leo	male	29.0	0	0	315082	

	fare	cabin	embarked	boat	body	home.dest
1308	7.875	NaN	S	NaN	NaN	NaN

3.6.4 See random N rows using .sample(N)

```
In [9]: titanic.sample(3)
```

```
Out[9]:
```

	pclass	survived	name	sex	age	\
811	3	0	Ford, Mrs. Edward (Margaret Ann Watson)	female	48.0	
348	2	0	Bracken, Mr. James H	male	27.0	
1075	3	0	Odahl, Mr. Nils Martin	male	23.0	

	sibsp	parch	ticket	fare	cabin	embarked	boat	body	\
811	1	3	W./C. 6608	34.375	NaN	S	NaN	NaN	
348	0	0	220367	13.000	NaN	S	NaN	NaN	
1075	0	0	7267	9.225	NaN	S	NaN	NaN	

	home.dest
811	Rotherfield, Sussex, England Essex Co, MA
348	Lake Arthur, Chavez County, NM
1075	NaN

3.7 len() return number of observations (rows)

```
In [10]: len(titanic)
```

```
Out[10]: 1309
```

3.8 .shape attribute gives the shape

```
In [11]: titanic.shape
```

```
Out[11]: (1309, 14)
```

3.9 .describe(): Get basic statistics across all columns

- Detects which columns are quantitative gives descriptive stats for those

```
In [12]: titanic.describe()
```

```
Out [12]:
```

	pclass	survived	age	sibsp	parch \
count	1309.000000	1309.000000	1046.000000	1309.000000	1309.000000
mean	2.294882	0.381971	29.881135	0.498854	0.385027
std	0.837836	0.486055	14.413500	1.041658	0.865560
min	1.000000	0.000000	0.166700	0.000000	0.000000
25%	2.000000	0.000000	21.000000	0.000000	0.000000
50%	3.000000	0.000000	28.000000	0.000000	0.000000
75%	3.000000	1.000000	39.000000	1.000000	0.000000
max	3.000000	1.000000	80.000000	8.000000	9.000000

	fare	body
count	1308.000000	121.000000
mean	33.295479	160.809917
std	51.758668	97.696922
min	0.000000	1.000000
25%	7.895800	72.000000
50%	14.454200	155.000000
75%	31.275000	256.000000
max	512.329200	328.000000

3.10 .count() give number of non-empty cells

```
In [13]: titanic.count()
```

```
Out [13]: pclass      1309
survived    1309
name        1309
sex         1309
age         1046
sibsp       1309
parch       1309
ticket      1309
fare        1308
cabin       295
embarked    1307
boat        486
body        121
home.dest   745
dtype: int64
```

3.11 DataFrame row and column headers

- Like a NumPy array, but with column and row headers.
- Enables slicing by headers, and not just indices like with NumPy arrays
- The collection of row headers is stored in the .index attribute.
- The collection of column headers is stored in the .columns attribute.

```
In [14]: titanic.columns
```

```
Out[14]: Index(['pclass', 'survived', 'name', 'sex', 'age', 'sibsp', 'parch', 'ticket',
               'fare', 'cabin', 'embarked', 'boat', 'body', 'home.dest'],
               dtype='object')
```

```
In [15]: titanic.index
```

```
Out[15]: RangeIndex(start=0, stop=1309, step=1)
```

3.12 Get a single column

Two ways to do it:

1. Use the “object-oriented” style of [API](#), i.e., the “dot.”
2. Use the dict style, i.e., key-value style (put the column name into brackets, get the column)
3. The returned data type is a PANDAS Series object, which keeps the index from the DataFrame attached

```
In [16]: titanic.columns
```

```
Out[16]: Index(['pclass', 'survived', 'name', 'sex', 'age', 'sibsp', 'parch', 'ticket',
               'fare', 'cabin', 'embarked', 'boat', 'body', 'home.dest'],
               dtype='object')
```

```
In [17]: titanic['home.dest']
```

```
Out[17]: 0                St Louis, MO
1    Montreal, PQ / Chesterville, ON
2    Montreal, PQ / Chesterville, ON
3    Montreal, PQ / Chesterville, ON
4    Montreal, PQ / Chesterville, ON
5                New York, NY
6                Hudson, NY
7                Belfast, NI
8    Bayside, Queens, NY
9    Montevideo, Uruguay
10               New York, NY
11               New York, NY
12               Paris, France
13                NaN
14    Hessle, Yorks
15               New York, NY
16               Montreal, PQ
17               Montreal, PQ
18                NaN
19    Winnipeg, MN
20               New York, NY
21               New York, NY
22               New York, NY
23                NaN
```

```

24                                     NaN
...
1284                                 NaN
1285                                 NaN
1286                                 NaN
1287                                 NaN
1288                                 NaN
1289                                 NaN
1290                                 NaN
1291                                 NaN
1292                                 NaN
1293                                 NaN
1294                                 NaN
1295                                 NaN
1296                                 NaN
1297                                 NaN
1298                                 NaN
1299                                 NaN
1300                                 NaN
1301                                 NaN
1302                                 NaN
1303                                 NaN
1304                                 NaN
1305                                 NaN
1306                                 NaN
1307                                 NaN
1308                                 NaN
Name: home.dest, Length: 1309, dtype: object

```

3.13 using .values

```
In [18]: titanic['home.dest'].values
```

```
Out[18]: array(['St Louis, MO', 'Montreal, PQ / Chesterville, ON',
               'Montreal, PQ / Chesterville, ON', ..., nan, nan, nan],
            dtype=object)
```

3.14 .value_counts()

```
In [19]: titanic['sex']
```

```
Out[19]: 0      female
         1      male
         2      female
         3      male
         4      female
         5      male
         6      female
```

```

7         male
8        female
9         male
10        male
11        female
12        female
13        female
14         male
15         male
16         male
17        female
18        female
19         male
20         male
21        female
22         male
23        female
24        female
...
1284       male
1285       male
1286      female
1287       male
1288       male
1289       male
1290      female
1291       male
1292       male
1293       male
1294       male
1295       male
1296       male
1297       male
1298       male
1299       male
1300      female
1301       male
1302       male
1303       male
1304      female
1305      female
1306       male
1307       male
1308       male
Name: sex, Length: 1309, dtype: object

```

```
In [20]: titanic.sex.value_counts()
```

```
Out[20]: male      843
```

```
female      466
Name: sex, dtype: int64
```

3.15 Use .pivot_table() to have a breakdown of the data

3.15.1 For categorical data, use aggfunc='count'

```
In [ ]: titanic.pivot_table?
```

```
In [21]: titanic.count()
```

```
Out[21]: pclass      1309
survived      1309
name          1309
sex           1309
age           1046
sibsp         1309
parch         1309
ticket        1309
fare          1308
cabin         295
embarked      1307
boat          486
body          121
home.dest     745
dtype: int64
```

```
In [22]: titanic.pivot_table( values='survived', index='pclass',
                               columns='sex', aggfunc='count',
                               margins=True)
```

```
Out[22]: sex      female  male  All
pclass
1          144    179    323
2          106    171    277
3          216    493    709
All         466    843   1309
```

3.15.2 For non-categorical data, can use another statistical measure for aggregation, like mean

```
In [23]: titanic.pivot_table( values='age', index='sex',
                               columns='pclass',
                               aggfunc='mean', margins=True)
```

```
Out[23]: pclass      1          2          3          All
sex
female  37.037594  27.499191  22.185307  28.687071
male    41.029250  30.815401  25.962273  30.585233
All     39.159918  29.506705  24.816367  29.881135
```


3.16 Quick figures

- Execute this Jupyter command `%matplotlib inline` before executing code that makes figures to get Jupyter to render them as output.

```
In [27]: %matplotlib inline
```

3.16.1 Univariate histograms

```
In [ ]: titanic.age.hist?
```

```
In [24]: thing = titanic.age
```


```
In [25]: type( thing)
```

```
Out[25]: pandas.core.series.Series
```

```
In [ ]: thing.hist?
```

```
In [28]: titanic.age.hist()
```


```
Out[28]: <matplotlib.axes._subplots.AxesSubplot at 0x11a320588>
```



3.16.2 Bivariate scatter plot using the .plot attribute

```
In [29]: titanic.plot.scatter( 'age', 'fare' )
```

```
Out[29]: <matplotlib.axes._subplots.AxesSubplot at 0x11a3a17b8>
```



3.17 Missing data in PANDAS

- Represented as `np.nan`, which stands for “Not A Number”
- NaN has type float
- No missing data representation for an integer!
 - Either convert all to floats to use NaN (recommended!), or
 - Convert values into strings and store empties as "" (less recommended)
 - Establish a “flag” value, e.g., -999 and filter out those before using (not recommended!)

In [30]: `import numpy as np`

In [31]: `np.nan`

Out[31]: `nan`

In [32]: `type(np.nan)`

Out[32]: `float`

3.18 Column data types

- A single column of data within a PANDAS DataFrame is called a Series.
- All values within a Series must be of the same type.
- Use the `.dtypes` attribute to check data types for each column

```
In [33]: titanic.head(3)
```

```
Out[33]:
```

	pclass	survived	name	sex	age	sibsp	\
0	1	1	Allen, Miss. Elisabeth Walton	female	29.0000	0	
1	1	1	Allison, Master. Hudson Trevor	male	0.9167	1	
2	1	0	Allison, Miss. Helen Loraine	female	2.0000	1	

	parch	ticket	fare	cabin	embarked	boat	body	\
0	0	24160	211.3375	B5	S	2	NaN	
1	2	113781	151.5500	C22 C26	S	11	NaN	
2	2	113781	151.5500	C22 C26	S	NaN	NaN	

	home.dest
0	St Louis, MO
1	Montreal, PQ / Chesterville, ON
2	Montreal, PQ / Chesterville, ON

```
In [34]: titanic.count()
```

```
Out[34]:
```

pclass	1309
survived	1309
name	1309
sex	1309
age	1046
sibsp	1309
parch	1309
ticket	1309
fare	1308
cabin	295
embarked	1307
boat	486
body	121
home.dest	745
dtype:	int64

```
In [35]: titanic.dtypes
```

```
Out[35]:
```

pclass	int64
survived	int64
name	object
sex	object
age	float64
sibsp	int64
parch	int64
ticket	object
fare	float64
cabin	object
embarked	object
boat	object

```
body          float64
home.dest     object
dtype: object
```

3.19 Column data types may hint at missing values

When using `pd.read_csv()` and `pd.read_excel()` to load a file from disk, PANDAS will try to pick a data type for a column that makes sense.

- If a float64 (just a fancy float), then missing values in the form of NaN are possible
 - Use `.count()` to count non-empty (non-NaN) values
- If an int64 (just a fancy int), then probably no missing values in that column
- If an object, this almost always means it's a string in there
 - Can represent missing values as "", but `.count()` only works for float data types!

```
In [36]: some_emptyys = pd.Series( ["", "asdf", "", "", "", "27", ""] )
        print( some_emptyys.dtype )
        some_emptyys.count()
```

```
object
```

```
Out[36]: 7
```

3.19.1 Coerce to numeric values using `pd.to_numeric()`

```
In [37]: some_emptyys = pd.to_numeric( some_emptyys, errors='coerce' )
```

```
In [38]: some_emptyys
```

```
Out[38]: 0      NaN
         1      NaN
         2      NaN
         3      NaN
         4      NaN
         5    27.0
         6      NaN
         dtype: float64
```

```
In [39]: print( some_emptyys.dtype )
        some_emptyys.count()
```

```
float64
```

```
Out[39]: 1
```

3.20 Statistics on a DataFrame ignore NaNs (as one might expect)

- In other words, doesn't count missing values as 0

```
In [40]: titanic.count()
```

```
Out[40]: pclass      1309
         survived    1309
         name        1309
         sex         1309
         age         1046
         sibsp       1309
         parch       1309
         ticket      1309
         fare        1308
         cabin       295
         embarked    1307
         boat        486
         body        121
         home.dest    745
         dtype: int64
```

```
In [41]: titanic.age.describe()
```

```
Out[41]: count      1046.000000
         mean        29.881135
         std         14.413500
         min         0.166700
         25%         21.000000
         50%         28.000000
         75%         39.000000
         max         80.000000
         Name: age, dtype: float64
```

3.21 Using the Seaborn Package for visualization

- Browse [this page](#) to see all the types of nice figures you can make


```
In [42]: import seaborn as sns
```

```
In [43]: import matplotlib.pyplot as plt
```

```
In [44]: fig, ax = plt.subplots( figsize=(5,4) )
```

```
         sns.swarmplot( x='sex', y='age', hue='survived',
                        data=titanic, ax=ax )
         #fig.savefig( 'testytest.pdf' )
```

```
Out[44]: <matplotlib.axes._subplots.AxesSubplot at 0x11a49b860>
```



```
In [45]: type( fig )
```

```
Out[45]: matplotlib.figure.Figure
```

```
In [46]: type( ax)
```


```
Out[46]: matplotlib.axes._subplots.AxesSubplot
```

```
In [ ]: ax.
```

```
In [ ]: sns.swarmplot?
```

```
In [47]: sns.swarmplot( x='pclass', y='age', hue='survived',  
                        data=titanic )
```

```
Out[47]: <matplotlib.axes._subplots.AxesSubplot at 0x1262ae240>
```



3.22 Subselecting based on one of the variables

```
In [48]: titanic.shape
```

```
Out[48]: (1309, 14)
```

```
In [49]: titanic.sex.value_counts()
```

```
Out[49]: male      843
         female    466
         Name: sex, dtype: int64
```

```
In [50]: titanic.sex.head()
```

```
Out[50]: 0    female
         1     male
         2    female
         3     male
         4    female
         Name: sex, dtype: object
```

```
In [51]: titanic.sex == 'male'
```

```
Out[51]: 0    False
         1     True
```

2	False
3	True
4	False
5	True
6	False
7	True
8	False
9	True
10	True
11	False
12	False
13	False
14	True
15	True
16	True
17	False
18	False
19	True
20	True
21	False
22	True
23	False
24	False
	...
1284	True
1285	True
1286	False
1287	True
1288	True
1289	True
1290	False
1291	True
1292	True
1293	True
1294	True
1295	True
1296	True
1297	True
1298	True
1299	True
1300	False
1301	True
1302	True
1303	True
1304	False
1305	False
1306	True
1307	True


```

1308      True
      Name: sex, Length: 1309, dtype: bool

In [52]: bool_array = titanic.sex == 'male'

In [53]: len(bool_array)

Out[53]: 1309

In [54]: (titanic.sex == 'male').head()

Out[54]: 0    False
         1     True
         2    False
         3     True
         4    False
         Name: sex, dtype: bool

In [55]: male = titanic[ titanic.sex == 'male' ]

In [56]: # Boolean selector array have to be the same shape as the array itself!!
         bool_array = [True]*1000

In [57]: #titanic[ bool_array ]

In [58]: male.shape

Out[58]: (843, 14)

In [59]: gender_tf = titanic.sex == 'male'

In [60]: gender_tf.shape

Out[60]: (1309,)

In [61]: male.shape

Out[61]: (843, 14)

In [62]: female = titanic[ titanic.sex == 'female' ]


In [63]: female.shape

Out[63]: (466, 14)

In [64]: sns.swarmplot( x='pclass', y='age', hue='survived',
                        data=male)


Out[64]: <matplotlib.axes._subplots.AxesSubplot at 0x12636ba20>

```



```
In [65]: sns.swarmplot( x='pclass', y='age', hue='survived',  
                        data=female )
```

```
Out[65]: <matplotlib.axes._subplots.AxesSubplot at 0x126370160>
```



3.23 Slicing by rows and columns using .loc[]

```
In [66]: subset = titanic[ titanic.age < 25 ]  
In [67]: subset.shape  
Out[67]: (409, 14)  
In [68]: subset = titanic.loc[ titanic.age < 25 ]  
In [69]: subset.shape  
Out[69]: (409, 14)
```

4 Complex sort

```
In [70]: age_bool = titanic.age < 10  
In [71]: age_bool.value_counts()  
Out[71]: False    1227  
         True      82  
         Name: age, dtype: int64  
In [72]: class_bool = titanic.pclass == 1  
In [73]: class_bool.value_counts()  
Out[73]: False    986  
         True    323  
         Name: pclass, dtype: int64  
In [74]: age_class_bool = age_bool & class_bool  
In [75]: age_class_bool.value_counts()  
Out[75]: False    1305  
         True      4  
         dtype: int64  
In [76]: titanic.loc[ age_class_bool, 'age' ]  
Out[76]: 1      0.9167  
         2      2.0000  
         94     4.0000  
         273    6.0000  
         Name: age, dtype: float64  
In [77]: len(subset)  
Out[77]: 409
```

4.1 Using .sort_values() for simple or complex sorting

```
In [ ]: titanic.sort_values?
```

```
In [78]: titanic.shape
```

```
Out[78]: (1309, 14)
```

```
In [82]: titanic.sort_values( by=['pclass', 'age']).head()
```

```
Out[82]:
```

	pclass	survived	name	sex	age	\
1	1	1	Allison, Master. Hudson Trevor	male	0.9167	
2	1	0	Allison, Miss. Helen Loraine	female	2.0000	
94	1	1	Dodge, Master. Washington	male	4.0000	
273	1	1	Spedden, Master. Robert Douglas	male	6.0000	
54	1	1	Carter, Master. William Thornton II	male	11.0000	

	sibsp	parch	ticket	fare	cabin	embarked	boat	body	\
1	1	2	113781	151.5500	C22 C26	S	11	NaN	
2	1	2	113781	151.5500	C22 C26	S	NaN	NaN	
94	0	2	33638	81.8583	A34	S	5	NaN	
273	0	2	16966	134.5000	E34	C	3	NaN	
54	1	2	113760	120.0000	B96 B98	S	4	NaN	

	home.dest
1	Montreal, PQ / Chesterville, ON
2	Montreal, PQ / Chesterville, ON
94	San Francisco, CA
273	Tuxedo Park, NY
54	Bryn Mawr, PA

```
In [83]: titanic.sort_values( by=['pclass', 'age'],  
                             ascending=False).head()
```

```
Out[83]:
```

	pclass	survived	name	sex	age	sibsp	parch	\
1235	3	0	Svensson, Mr. Johan	male	74.0	0	0	
727	3	0	Connors, Mr. Patrick	male	70.5	0	0	
782	3	0	Duane, Mr. Frank	male	65.0	0	0	
1261	3	1	Turkula, Mrs. (Hedwig)	female	63.0	0	0	
1068	3	0	Nysveen, Mr. Johan Hansen	male	61.0	0	0	

	ticket	fare	cabin	embarked	boat	body	home.dest
1235	347060	7.7750	NaN	S	NaN	NaN	NaN
727	370369	7.7500	NaN	Q	NaN	171.0	NaN
782	336439	7.7500	NaN	Q	NaN	NaN	NaN
1261	4134	9.5875	NaN	S	15	NaN	NaN
1068	345364	6.2375	NaN	S	NaN	NaN	NaN

```
In [84]: titanic['home.dest'].sample(10)
```

```
Out[84]: 568      NaN
        1147      NaN
        786      Tofta, Sweden Joliet, IL
        1104      NaN
        604      Norway Los Angeles, CA
        94       San Francisco, CA
        784      West Haven, CT
        125      NaN
        864      NaN
        214      Lexington, MA
        Name: home.dest, dtype: object
```