

NIAPythonDay4

March 28, 2019

1 Day 1: IDE, basic data types, operators

2 Day 2: Slicing, NumPy array

3 Day 3: Pandas DataFrame

4 Day 4: Case Study: Microarray Data analysis and Visualization

- file: ExampleMicroarrayData.xls
- Data: brain region microarray data (measures gene expression levels in a tissue)
- Experimental variables = tissue type
 - hippocampus & cerebral cortex
- Columns: 4 repeats for each tissue type
- Rows: 1 row per read; 1 or more rows per gene. Redundant reads need to be merged.

4.1 Analysis Method

1. Import file
2. For a given gene with redundant reads, take maximum value (collapse multiple rows into one row)
3. Log transform
4. Z-score
5. For each comparison, take difference of means of repeats and p-value
6. Create ribbon plot for genes with fold change of 1 and p-value < 0.05

4.2 Import libraries

```
In [1]: import pandas as pd
import numpy as np
```

4.3 Optional: set Pandas display precision

Use TAB key to show you what options are available

```
In [ ]: pd.options.display
```

```
In [2]: pd.options.display.precision = 3
```

4.4 Read in Excel File

```
In [3]: microa = pd.read_excel( 'ExampleMicroarrayData.xls')
```

4.5 See what we got

```
In [4]: microa.head()
```

```
Out [4]:
```

	ArrayID	Symbol	BR1_1	BR1_2	BR1_3	BR1_4	BR2_1	\
0	1	NA1	24540.450	43972.090	32894.380	39717.850	37762.740	
1	2	NA2	3.384	3.420	3.184	3.504	2.218	
2	3	NA3	3.402	3.452	3.215	3.549	2.244	
3	4	Tbc1d19	497.424	704.163	598.064	770.349	318.999	
4	5	Cfc1	3.436	14.535	3.271	3.640	2.296	

		BR2_2	BR2_3	BR2_4
0	58349.680	20307.620	25211.600	
1	2.300	2.404	2.307	
2	2.327	2.434	2.334	
3	315.689	329.867	342.941	
4	2.374	10.220	2.384	

```
In [5]: microa.shape
```

```
Out [5]: (59734, 10)
```

```
In [6]: microa.columns
```

```
Out [6]: Index(['ArrayID', 'Symbol', 'BR1_1', 'BR1_2', 'BR1_3', 'BR1_4', 'BR2_1',  
               'BR2_2', 'BR2_3', 'BR2_4'],  
              dtype='object')
```

4.6 Review: Get basic statistics across all columns using .describe()

```
In [7]: microa.describe()
```

```
Out [7]:
```

	ArrayID	BR1_1	BR1_2	BR1_3	BR1_4	BR2_1	\
count	59734.000	59734.000	59734.000	59734.000	59734.000	59734.000	
mean	31524.803	3826.969	4417.788	4440.600	5140.971	2811.512	
std	18161.786	12949.801	14653.350	14482.841	16651.936	11257.903	
min	1.000	3.126	2.979	2.860	3.060	2.048	
25%	15839.250	18.760	22.889	21.042	24.194	13.127	
50%	31507.500	292.898	361.491	363.980	418.769	205.459	
75%	47259.750	2143.263	2576.601	2603.771	3029.693	1387.381	
max	62976.000	309771.300	351175.700	340936.200	377660.600	548163.100	

	BR2_2	BR2_3	BR2_4
count	59734.000	59734.000	59734.000
mean	3070.629	3017.429	3422.667

std	11271.983	11149.905	12554.989
min	2.024	2.049	2.084
25%	13.279	10.264	12.211
50%	226.891	187.594	224.595
75%	1625.061	1499.185	1726.906
max	327943.800	381527.600	343815.500

4.7 Review: subselect rows by boolean criterion using brackets []

```
In [8]: microa[ microa.Symbol == 'Mfsd9' ]
```

```
Out [8]:
```

	ArrayID	Symbol	BR1_1	BR1_2	BR1_3	BR1_4	BR2_1	\
3692	3920	Mfsd9	968.993	1010.708	1007.123	1130.436	434.947	
14504	15379	Mfsd9	949.827	1056.157	979.577	1163.150	482.391	
15412	16338	Mfsd9	23.193	18.875	22.034	24.326	8.435	
19317	20435	Mfsd9	891.462	1036.658	991.143	1173.474	452.980	
21174	22376	Mfsd9	920.478	998.698	1052.486	1231.313	434.002	
26077	27531	Mfsd9	18.640	24.579	12.145	28.158	9.539	
32783	34582	Mfsd9	854.077	1062.900	1008.647	1180.013	435.248	
33005	34823	Mfsd9	949.943	985.132	1079.542	1201.548	463.207	
34495	36380	Mfsd9	856.771	1024.499	914.007	1123.651	473.591	
48207	50847	Mfsd9	955.400	1029.965	1039.952	1146.253	471.125	
57835	60983	Mfsd9	915.912	1021.728	947.302	1168.269	421.222	
58099	61263	Mfsd9	1051.273	1004.055	996.786	1133.530	448.804	

	BR2_2	BR2_3	BR2_4
3692	589.708	600.502	621.537
14504	543.322	601.314	674.765
15412	9.239	7.162	9.467
19317	555.692	622.570	606.833
21174	547.448	637.982	625.079
26077	12.697	8.741	11.155
32783	577.290	663.561	615.494
33005	530.134	596.087	659.742
34495	585.492	618.376	644.952
48207	596.579	627.870	621.421
57835	589.248	632.728	660.304
58099	589.341	651.024	605.704

4.8 Review: subselect one or more columns using brackets []

```
In [ ]: microa[ 'Symbol' ]
```

```
In [ ]: microa[ ['ArrayID', 'Symbol' ] ]
```

4.9 For subselecting rows and columns at the same time, use .loc[] (or.iloc[])

```
In [24]: microa.loc[ microa.Symbol == 'Mfsd9',
                    ['BR1_1', 'BR2_1' ] ]
```

```
Out [24]:
```

	BR1_1	BR2_1
3692	968.993	434.947
14504	949.827	482.391
15412	23.193	8.435
19317	891.462	452.980
21174	920.478	434.002
26077	18.640	9.539
32783	854.077	435.248
33005	949.943	463.207
34495	856.771	473.591
48207	955.400	471.125
57835	915.912	421.222
58099	1051.273	448.804

4.10 Note: Difference between .loc[] and .iloc[]

- Use .loc[] to slice by row/column *NAMES* or *BOOLEANS*
- Use .iloc[] to slice by row/column *INDICES* (Like we used with NumPy)

```
In [ ]: microa.loc[ :10, -4: ]
```

```
In [26]: microa.iloc[ :10, -4: ]
```

```
Out [26]:
```

	BR2_1	BR2_2	BR2_3	BR2_4
0	37762.740	58349.680	20307.620	25211.600
1	2.218	2.300	2.404	2.307
2	2.244	2.327	2.434	2.334
3	318.999	315.689	329.867	342.941
4	2.296	2.374	10.220	2.384
5	60.898	79.251	88.181	83.621
6	337.046	246.692	264.554	373.699
7	14415.300	15068.380	15317.930	18569.540
8	21151.840	23888.320	25168.000	28457.830
9	266.326	225.108	254.067	244.691

4.11 How many unique gene symbols do we have?

```
In [27]: len( microa.Symbol )
```

```
Out [27]: 59734
```

```
In [28]: microa.Symbol.unique()
```

```
Out [28]: array(['NA1', 'NA2', 'NA3', ..., 'NA11372', 'NA11373', 'NA11374'],
                dtype=object)
```

```
In [29]: len(microa.Symbol.unique())
```

```
Out [29]: 29784
```

```

In [ ]: microa.Symbol.value_counts()

In [31]: microa.Symbol.value_counts().value_counts()

Out[31]: 1      14855
         2      11361
         4       1162
        10       948
         3       932
        12       211
        11       130
         5       117
         6        28
        13        16
        20         8
         7         7
        14         3
         9         3
         8         2
        16         1
        Name: Symbol, dtype: int64

```

4.12 Group rows by gene symbol using .groupby()

- Pass the function the name of the column containing the values you want to group by

```

In [32]: grouped = microa.groupby('Symbol')

In [33]: type( grouped)

Out[33]: pandas.core.groupby.groupby.DataFrameGroupBy

In [ ]: grouped.mean()

```

4.12.1 Get a certain group using .get_group()

```

In [39]: grouped.get_group( 'Mfsd9' )

Out[39]:

```

	ArrayID	BR1_1	BR1_2	BR1_3	BR1_4	BR2_1	BR2_2	\
3692	3920	968.993	1010.708	1007.123	1130.436	434.947	589.708	
14504	15379	949.827	1056.157	979.577	1163.150	482.391	543.322	
15412	16338	23.193	18.875	22.034	24.326	8.435	9.239	
19317	20435	891.462	1036.658	991.143	1173.474	452.980	555.692	
21174	22376	920.478	998.698	1052.486	1231.313	434.002	547.448	
26077	27531	18.640	24.579	12.145	28.158	9.539	12.697	
32783	34582	854.077	1062.900	1008.647	1180.013	435.248	577.290	
33005	34823	949.943	985.132	1079.542	1201.548	463.207	530.134	
34495	36380	856.771	1024.499	914.007	1123.651	473.591	585.492	
48207	50847	955.400	1029.965	1039.952	1146.253	471.125	596.579	

57835	60983	915.912	1021.728	947.302	1168.269	421.222	589.248
58099	61263	1051.273	1004.055	996.786	1133.530	448.804	589.341

	BR2_3	BR2_4
3692	600.502	621.537
14504	601.314	674.765
15412	7.162	9.467
19317	622.570	606.833
21174	637.982	625.079
26077	8.741	11.155
32783	663.561	615.494
33005	596.087	659.742
34495	618.376	644.952
48207	627.870	621.421
57835	632.728	660.304
58099	651.024	605.704

```
In [40]: max_expression = microa.groupby('Symbol').max()
```

```
In [41]: max_expression.shape
```

```
Out[41]: (29784, 9)
```

```
In [44]: max_expression.head()
```

```
Out[44]:
```

	ArrayID	BR1_1	BR1_2	BR1_3	BR1_4	BR2_1	BR2_2	\
Symbol								
A1bg	47917	6.551	9.730	6.253	13.327	3.550	2.836	
A1cf	50995	3.718	3.788	3.795	5.505	2.599	2.552	
A1i3	59848	3.740	7.137	3.697	4.011	2.605	2.541	
A26c2	52963	3.738	3.796	3.797	11.602	5.992	3.765	
A2m	36219	6142.893	5840.628	6192.190	6855.576	533.450	408.658	

	BR2_3	BR2_4
Symbol		
A1bg	4.228	4.405
A1cf	2.595	2.649
A1i3	2.601	2.648
A26c2	2.597	8.467
A2m	355.236	371.571

```
In [43]: max_expression.iloc[ 4, 1]
```

```
Out[43]: 6142.893
```

```
In [ ]: max_expression.loc[ 'A2m', 'BR1_1']
```

4.13 Combine reads

Take maximum expression level for a given gene.

```
In [45]: max_expression = grouped.max()
```

```
In [46]: max_expression.shape
```

```
Out[46]: (29784, 9)
```

5 Drop the Array ID Column

```
In [47]: max_expression.head()
```

```
Out[47]:
```

	ArrayID	BR1_1	BR1_2	BR1_3	BR1_4	BR2_1	BR2_2	\
Symbol								
A1bg	47917	6.551	9.730	6.253	13.327	3.550	2.836	
A1cf	50995	3.718	3.788	3.795	5.505	2.599	2.552	
A1i3	59848	3.740	7.137	3.697	4.011	2.605	2.541	
A26c2	52963	3.738	3.796	3.797	11.602	5.992	3.765	
A2m	36219	6142.893	5840.628	6192.190	6855.576	533.450	408.658	

	BR2_3	BR2_4
Symbol		
A1bg	4.228	4.405
A1cf	2.595	2.649
A1i3	2.601	2.648
A26c2	2.597	8.467
A2m	355.236	371.571

```
In [ ]: max_expression.drop?
```

```
In [48]: max_expression.drop( columns='ArrayID', inplace=True )
```

```
In [49]: max_expression.head()
```

```
Out[49]:
```

	BR1_1	BR1_2	BR1_3	BR1_4	BR2_1	BR2_2	BR2_3	\
Symbol								
A1bg	6.551	9.730	6.253	13.327	3.550	2.836	4.228	
A1cf	3.718	3.788	3.795	5.505	2.599	2.552	2.595	
A1i3	3.740	7.137	3.697	4.011	2.605	2.541	2.601	
A26c2	3.738	3.796	3.797	11.602	5.992	3.765	2.597	
A2m	6142.893	5840.628	6192.190	6855.576	533.450	408.658	355.236	

	BR2_4
Symbol	
A1bg	4.405
A1cf	2.649
A1i3	2.648
A26c2	8.467
A2m	371.571

5.1 .apply() a log transformation function over all expression values

```
In [ ]: np.log2?
In [51]: np.log2( [1, 10, 100, 1000] )
Out[51]: array([0.          , 3.32192809, 6.64385619, 9.96578428])
In [52]: max_expression.shape
Out[52]: (29784, 8)
In [53]: log_trans = max_expression.apply( np.log2 )
In [ ]: log_trans
```

5.2 .apply() a Z-score transformation function over all genes

```
In [55]: from scipy.stats import zscore
In [ ]: zscore?
In [57]: zscore( [ 1,2,3,4,5,6,7,8,9,10] )
Out[57]: array([-1.5666989 , -1.21854359, -0.87038828, -0.52223297, -0.17407766,
                0.17407766,  0.52223297,  0.87038828,  1.21854359,  1.5666989 ])
In [58]: z_trans = log_trans.apply(zscore)
In [60]: log_trans.head()
Out[60]:
```


	BR1_1	BR1_2	BR1_3	BR1_4	BR2_1	BR2_2	BR2_3	BR2_4
Symbol								
A1bg	2.712	3.283	2.644	3.736	1.828	1.504	2.080	2.139
A1cf	1.895	1.921	1.924	2.461	1.378	1.352	1.376	1.405
A1i3	1.903	2.835	1.886	2.004	1.381	1.345	1.379	1.405
A26c2	1.902	1.924	1.925	3.536	2.583	1.913	1.377	3.082
A2m	12.585	12.512	12.596	12.743	9.059	8.675	8.473	8.537

```
In [ ]: log_trans.apply?
In [61]: log_trans.shape
Out[61]: (29784, 8)
In [62]: z_trans.shape
Out[62]: (29784, 8)
In [63]: z_trans.head()
Out[63]:
```

	BR1_1	BR1_2	BR1_3	BR1_4	BR2_1	BR2_2	BR2_3	BR2_4
Symbol								
A1bg	-1.104	-1.023	-1.147	-0.921	-1.198	-1.271	-1.085	-1.105
A1cf	-1.301	-1.350	-1.317	-1.219	-1.307	-1.307	-1.250	-1.276
A1i3	-1.299	-1.131	-1.326	-1.326	-1.306	-1.308	-1.250	-1.276
A26c2	-1.299	-1.349	-1.317	-0.968	-1.017	-1.174	-1.250	-0.886
A2m	1.273	1.188	1.203	1.185	0.541	0.422	0.413	0.383

5.3 Compare: Make a histogram of fold change

```
In [64]: br1 = z_trans[ ['BR1_1', 'BR1_2', 'BR1_3', 'BR1_4' ] ]
In [65]: br1.shape
Out[65]: (29784, 4)
In [66]: br2 = z_trans[ ['BR2_1', 'BR2_2', 'BR2_3', 'BR2_4' ] ]
In [67]: br2.shape
Out[67]: (29784, 4)
In [68]: br1_mean = br1.mean(axis=1)
          br2_mean = br2.mean(axis=1)
In [69]: br2_mean.shape
Out[69]: (29784,)
In [70]: diff = br2_mean - br1_mean
In [71]: diff.shape
Out[71]: (29784,)
In [72]: # tell Jupyter Notebook to show the figure after it made it
          %matplotlib inline
In [73]: diff.hist()
Out[73]: <matplotlib.axes._subplots.AxesSubplot at 0x112b8fac8>
```



5.4 Subselect genes based on having fold change > 1

```
In [ ]: diff[ diff.abs() > 1 ]
```

5.5 Get gene expression difference and p-value for desired comparisons

Use Wilcoxon rank-sum statistic for two samples to test if the reads for a given gene across comparison groups come from different distributions.

```
In [75]: from scipy.stats import ranksums
```

```
In [76]: ranksums( [1,2,3,4], [5,6,7,8])
```

```
Out[76]: RanksumsResult(statistic=-2.3094010767585034, pvalue=0.020921335337794014)
```

```
In [77]: ranksums( [1,2,3,4], [50,60,70,80])
```

```
Out[77]: RanksumsResult(statistic=-2.3094010767585034, pvalue=0.020921335337794014)
```

```
In [78]: ranksums( [1,2,3,5], [4,6,7,8])
```

```
Out[78]: RanksumsResult(statistic=-2.0207259421636903, pvalue=0.043308142810791955)
```

5.5.1 Go row-by-row doing significance test

```
In [79]: import time
```

```
In [80]: # create an empty list onto which we can append p-values
        pval_list = []
```

```
In [81]: t1 = time.time()
```

```
    # iterate over every gene
    for gene in br1.index:

        # subselect the expression values for the corresponding gene
        vals1 = br1.loc[ gene ]
        vals2 = br2.loc[ gene ]

        # Do the statistical test for this gene
        statistic, pvalue = ranksums( vals1, vals2 )

        # save the p-value
        pval_list.append( pvalue )

    t2 = time.time()
    print( "This operation took", t2-t1, "seconds.")
```

This operation took 13.183716058731079 seconds.

```
In [82]: len(pval_list)
```

```
Out[82]: 29784
```

```
In [83]: # One-liner!
```

```
t1 = time.time()
pval_list = [ ranksums(a,b)[1] for a, b in zip( br1.as_matrix(), br2.as_matrix() ) ]
t2 = time.time()
print( "This operation took", t2-t1, "seconds.")
```

```
/usr/local/lib/python3.6/site-packages/ipykernel_launcher.py:3: FutureWarning: Method .as_matrix() is deprecated,
This is separate from the ipykernel package so we can avoid doing imports until
```

This operation took 4.83489203453064 seconds.

5.6 Combine fold change and p-value into one DataFrame

- Problem 1: the diff object is just one column (a Pandas “Series” object), not a full-fledged DataFrame
- Problem 2: the p-values are stored in a simple Python list
- Solution: use the to_frame() function to turn the Series into a DataFrame, then add the p-values as a new column to that DataFrame.

```
In [84]: type( diff )
```

```
Out[84]: pandas.core.series.Series
```

```
In [85]: diff.head()
```

```
Out[85]: Symbol
A1bg      -0.116
A1cf       0.012
A1i3      -0.015
A26c2     0.151
A2m       -0.772
dtype: float64
```

```
In [86]: combined = diff.to_frame()
```

```
In [87]: combined.head()
```

```
Out[87]:          0
Symbol
A1bg      -0.116
A1cf       0.012
A1i3      -0.015
A26c2     0.151
A2m       -0.772
```

```
In [88]: combined = diff.to_frame( name='fold')
```

```
In [89]: combined.head()
```

```
Out[89]:
```

	fold
Symbol	
A1bg	-0.116
A1cf	0.012
A1i3	-0.015
A26c2	0.151
A2m	-0.772

```
In [90]: combined[ 'pvals' ] = pval_list
```

```
In [91]: combined.head()
```

```
Out[91]:
```

	fold	pvals
Symbol		
A1bg	-0.116	0.149
A1cf	0.012	0.564
A1i3	-0.015	0.564
A26c2	0.151	0.149
A2m	-0.772	0.021

```
In [92]: combined.shape
```

```
Out[92]: (29784, 2)
```

5.7 Subselect genes with fold change > 1

```
In [93]: combined[ combined.fold.abs() > 1 ]
```

```
Out[93]:
```

	fold	pvals
Symbol		
Adam33	1.231	0.021
C1ql2	-1.092	0.021
Cd3e	-1.243	0.021
Chst9	-1.406	0.021
Cox6a2	1.153	0.021
Csap1	-1.359	0.021
Cxcr1	-1.511	0.021
Cyp11b1	1.640	0.021
Dpp4	1.057	0.021
Fer1l4	1.212	0.021
Fibcd1	-1.098	0.021
Frem3	-1.127	0.021
Gpat2	1.229	0.021
Htr5b	-1.612	0.021
Klk8	-1.738	0.021

LOC688459	-1.045	0.021
Lhx9	-1.354	0.021
Lrrc10b	-1.367	0.021
Meox1	-1.040	0.021
Meox2	-1.028	0.021
NA1086	1.151	0.021
NA1528	1.069	0.083
NA3871	1.197	0.021
NA3959	1.206	0.021
NA4283	1.426	0.021
NA5561	-1.124	0.021
Ndst4	-1.346	0.021
Nhlh1	-1.090	0.021
Ntf3	-1.082	0.021
Ntrk1	-1.386	0.021
Nts	-2.095	0.021
Ocm2	1.288	0.021
Olr315	1.008	0.083
Olr59	1.267	0.021
Qrfpr	1.016	0.021
Slc9a4	-1.078	0.021
Sprr1a	1.146	0.021
Upb1	-1.026	0.021

5.8 Subselect genes with fold change > 1 AND p-value < 0.05

In [94]: `combined[(combined.fold.abs() > 1) & (combined.pvals < 0.05)]`

Out[94]:

	fold	pvals
Symbol		
Adam33	1.231	0.021
C1ql2	-1.092	0.021
Cd3e	-1.243	0.021
Chst9	-1.406	0.021
Cox6a2	1.153	0.021
Csap1	-1.359	0.021
Cxcr1	-1.511	0.021
Cyp11b1	1.640	0.021
Dpp4	1.057	0.021
Fer1l4	1.212	0.021
Fibcd1	-1.098	0.021
Frem3	-1.127	0.021
Gpat2	1.229	0.021
Htr5b	-1.612	0.021
Klk8	-1.738	0.021
LOC688459	-1.045	0.021
Lhx9	-1.354	0.021
Lrrc10b	-1.367	0.021

Meox1	-1.040	0.021
Meox2	-1.028	0.021
NA1086	1.151	0.021
NA3871	1.197	0.021
NA3959	1.206	0.021
NA4283	1.426	0.021
NA5561	-1.124	0.021
Ndst4	-1.346	0.021
Nhlh1	-1.090	0.021
Ntf3	-1.082	0.021
Ntrk1	-1.386	0.021
Nts	-2.095	0.021
Ocm2	1.288	0.021
Olr59	1.267	0.021
Qrfpr	1.016	0.021
Slc9a4	-1.078	0.021
Sprri1a	1.146	0.021
Upb1	-1.026	0.021

```
In [95]: plot_these = combined[ (combined.fold.abs() > 1) & (combined.pvals < 0.05) ]
```

```
In [96]: len(plot_these)
```

```
Out[96]: 36
```

Turn the row labels into variables in their own right:

```
In [97]: plot_these['Gene'] = plot_these.index
```

```
/usr/local/lib/python3.6/site-packages/ipykernel_launcher.py:1: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html>
 """Entry point for launching an IPython kernel.

```
In [98]: plot_these.sort_values(by='fold', inplace=True)
```

```
/usr/local/lib/python3.6/site-packages/ipykernel_launcher.py:1: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame
```

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html>
 """Entry point for launching an IPython kernel.

5.9 Generate RibbonPlot

5.9.1 Load some of Python's figure-making libraries:

```
In [99]: import seaborn as sns
import matplotlib.pyplot as plt
```

5.9.2 Set the style of the figure

When using the plotting package seaborn, there are five figure [styles](#) to choose from: 1. darkgrid 2. whitegrid 3. dark 4. white 5. ticks

```
In [100]: sns.set( style="whitegrid" )
```

5.9.3 Making the figure

```
In [101]: # I want this figure to be 6 inches wide and 10 inches tall
fig_dimensions=(6, 10)

# Create a blank figure to hang the data off of
figure, axes = plt.subplots( figsize=fig_dimensions )


# Plot the data onto the figure
sns.barplot( data=plot_these, x="fold", y="Gene" )

# Assign a title to the figure
chart_title = """Normalized difference in gene expression
Hippocampus vs. Cerebral Cortex"""
axes.set_title( chart_title, size=18 )

# Assign a label to the x-axis
axes.set_xlabel( "Fold change" )

Out[101]: Text(0.5,0,'Fold change')
```

Normalized difference in gene expression Hippocampus vs. Cerebral Cortex



5.9.4 Save the figure as a PDF:

```
In [ ]: figure.savefig( "ribbonplot.pdf")
```